Reverse-Engineering Perplexity AI's Visual Brand: A Modular Prompt System for AI Image Generation

Executive Summary

Perplexity AI has strategically cultivated a distinctive brand aesthetic, characterized by a "surreal-real with retro vibes" visual identity. This approach deliberately deviates from conventional sterile AI branding, aiming instead to convey intelligence, accessibility, and a human-centric approach to knowledge discovery. The visual style, deeply rooted in the brand's philosophy of providing clear, cited answers, subtly blends advanced capabilities with a reassuring sense of familiarity. This report meticulously deconstructs Perplexity's brand identity, analyzing its core philosophy, visual elements, and the nuanced interpretation of its unique aesthetic.

A comprehensive evaluation of leading AI image generation models—including GPT-Image-1, Midjourney v6, FLUX.1, and DALL-E 3—reveals their respective strengths and limitations in replicating this specific style. Based on this analysis, a modular, production-ready prompt system is proposed. This system leverages structured prompt components for style, composition, color, subject, and technical specifications, alongside negative prompts, to ensure reliable generation of adjacent visuals across diverse AI models. The modular architecture enhances consistency, reusability, and control, providing a scalable framework for maintaining brand visual integrity in dynamic creative workflows. Key recommendations include model-specific prompt adaptations, a structured workflow for iterative refinement, and continuous monitoring to uphold brand fidelity.

1. Deconstructing Perplexity AI's Visual Brand Identity

Perplexity AI's brand identity is a carefully constructed synthesis of its mission and aesthetic principles, designed to carve out a unique space in the competitive AI landscape. Understanding this identity requires examining its foundational philosophy and the deliberate choices made in its visual elements.

1.1 Perplexity's Core Brand Philosophy and Messaging

Perplexity AI positions itself as more than just another search tool; it is presented as "a fundamentally new way to seek knowledge". The company's messaging emphasizes being "intelligent, accessible, and alive," actively rejecting the "cold, hyper-minimalist designs" often associated with AI brands. This deliberate differentiation strategy aims to cultivate a brand that feels "human, curious, and open," prioritizing the questions as much as the answers.

The core mission of Perplexity is to provide "a direct line to the world's knowledge — compressed, cited, and made clear". This is underpinned by a commitment to "No gimmicks. No fluff. Just answers that make sense". This dedication to clarity and trustworthiness extends beyond its functional utility to its visual presentation. The brand's strategic market positioning is

deeply embedded in its human-centric approach. In an environment where AI is often perceived as robotic or impersonal, Perplexity consciously fosters an emotional connection with its audience. The "surreal-real" aspect of its visual style visually communicates the advanced, almost miraculous capability of AI in synthesizing vast amounts of information, while the "retro vibes" ground this advanced technology in familiarity and human experience. This combination effectively counters any perception of the brand being sterile or unapproachable, reinforcing its commitment to being a reliable and understandable source of knowledge. The self-definition as an "answer engine" serves as a profound philosophical anchor for its visuals. The brand's promise of delivering "answers that make sense" with clarity and directness implies a visual language that is both insightful and credible. The "surreal-real" component visually articulates the profound, almost magical access to expansive knowledge, while the "real" aspect ensures that this knowledge is presented in a verifiable and grounded manner, thereby reinforcing credibility. The "retro vibes" subtly evoke a sense of established wisdom or a classic, dependable source of information, offering a comforting contrast to the often overwhelming and chaotic nature of modern digital information. This indicates that the visuals should convey both cutting-edge capability and reassuring simplicity, aligning with the brand's overall promise.

1.2 Analysis of Key Visual Elements: Logo, Color Palette, and Typography

Perplexity's visual identity is meticulously crafted, with each element contributing to its distinctive brand persona.

The **logo** is described as a "brilliant fusion of meaning and simplicity". It appears as a symmetrical emblem but is composed of four symbolic elements: a cursor, representing the starting point of a search; an asterisk, symbolizing discovery and expansion; an open book, signifying knowledge, learning, and exploration; and intersecting windows, nodding to layered perspectives and new ways of seeing information. This design is not merely an abstract mark; it is a philosophical statement distilled into geometry, visually reinforcing Perplexity's mission to change how knowledge is accessed by expanding perspectives and creating connections. This semantic layering in the visuals creates a deeper connection with the user. The explicit design of the logo with these symbolic elements reveals a deep semantic intent behind Perplexity's visual identity. For AI image generation, this implies that prompts should not solely focus on aesthetic descriptions. Instead, they should aim to subtly embed conceptual meaning or narrative elements, even within abstract visuals. The "surreal-real" aspect can manifest as visuals that are not only aesthetically pleasing but also carry a deeper, perhaps symbolic, layer of meaning related to knowledge, discovery, or interconnectedness, fostering a more profound brand connection.

The **color palette** is "vibrant yet refined," built around "shades of teal, sky blue, and warm earth tones". These colors are intentionally chosen to convey "trust and curiosity". The palette is designed to feel "editorial and modern," drawing inspiration from "magazines and traditional media design". This choice positions Perplexity not just as an AI tool but as a trusted source of knowledge, where learning feels natural, not transactional. The "editorial modernity" serves as the sophisticated core of the "retro vibes." The explicit inspiration from "magazines and traditional media design" for the "editorial and modern" feel is a critical link to the "retro vibes" aspect of the brand. "Retro" here does not signify kitsch or overt nostalgia; instead, it harks back to the refined, curated aesthetic of traditional print media, which historically conveyed authority, trustworthiness, and intellectual depth. This "editorial" quality, particularly when combined with

"warm earth tones," creates a sense of established, reliable knowledge, contrasting sharply with the often chaotic and ephemeral nature of modern digital information. This suggests that AI-generated visuals should strive for a polished, thoughtful composition with a timeless quality, rather than relying on superficial vintage filters.

For **typography**, Perplexity uses FK Grotesk, described as the brand's "unspoken voice" that speaks in a "precise but approachable tone". This typeface features "sharp angles and subtle ink traps," giving it a "slightly mechanical feel without losing warmth". It is characterized as "rigid but not cold, structured but still inviting"—a perfect metaphor for the brand itself. This "structured but warm" paradox acts as a brand metaphor. The characterization of FK Grotesk typography perfectly encapsulates Perplexity's overall brand strategy. It embodies the duality of leveraging advanced AI (structured, precise, intelligent) while simultaneously maintaining a human, accessible, and approachable interface (warm, inviting). This duality should extend to the broader visual style. The "surreal-real" elements, for instance, should not be chaotic or random but rather structured and intentional, perhaps with clean lines and balanced compositions, even if the subject matter is abstract. Similarly, the "retro vibes" should be controlled and subtle, perhaps conveyed through precise color grading or compositional cues, rather than overwhelming stylistic effects. This implies a need for prompts that emphasize clarity, controlled complexity, and a harmonious blend of seemingly opposing attributes.

The following table summarizes these brand visual attributes and their interpretations for AI visual generation:

Table 1: Perplexity AI Brand Visual Attributes & Interpretations

Attribute Category	Specific Elements	Stated Meaning/Purpose	Interpretation for AI Visuals
Logo	Cursor	Starting point of a search	Subtle symbolic elements, guiding viewer's focus
	Asterisk	Discovery and expansion	Visual cues for new information, growth, or unexpected connections
	Open Book	Knowledge, learning, exploration	Representations of wisdom, data, or learning journeys
	Intersecting Windows	Layered perspectives, new ways of seeing information	Overlapping elements, depth, or visual metaphors for interconnectedness
Color Palette	Teal, Sky Blue	Trust and curiosity	Dominant cool tones for clarity and intelligence
	Warm Earth Tones	Trust and curiosity	Complementary warm tones for approachability and groundedness
	Overall	Editorial and modern, magazine-inspired	Refined, curated aesthetic; balanced and harmonious color schemes

Attribute Category	Specific Elements	Stated Meaning/Purpose	Interpretation for AI Visuals
Typography	FK Grotesk	Precise but approachable, rigid but not cold, structured but inviting	Clean, legible text integrated naturally; structured yet inviting compositions
Overall Tone	Minimal, Precise, Direct, Human, Curious, Open	Clear, confident, intellectual, accessible	Controlled complexity; harmonious blend of abstract and tangible elements; sophisticated rather than overtly playful

1.3 Interpreting the "Surreal-Real with Retro Vibes" Aesthetic

The "surreal-real with retro vibes" aesthetic is a sophisticated blend that defines Perplexity AI's visual communication. This section elaborates on the components of this unique style.

1.3.1 Elements of "Surreal-Real" in Perplexity's Visuals

The "surreal-real" aspect of Perplexity's visual identity is deeply connected to its mission of expanding perspectives and offering new ways of understanding information. This aesthetic aligns with the literary concept of "magical realism," where "the impossible coexists seamlessly with the ordinary". In this context, AI has normalized what might otherwise feel miraculous, such as instant access to vast, synthesized knowledge. Surrealistic techniques, including "combining unrelated elements alongside dream-based imagery and symbolic meanings, distortions and automatic artistic processes" within realistic media art, contribute to this effect.
The approach to surrealism here is conceptual rather than overtly fantastical, fostering intellectual engagement. The research highlights "layered perspectives" and the seamless coexistence of the impossible with the ordinary, as seen in magical realism. This indicates that the surreal element should not manifest as chaotic or unsettling imagery. Instead, it aims to present complex ideas or new ways of understanding information in a visually intriguing, almost dreamlike, yet fundamentally grounded manner. This implies that AI prompts should focus on subtle juxtapositions, symbolic elements, or unexpected perspectives within otherwise realistic scenes, rather than generating purely abstract or non-representational surrealism. The goal is to make the viewer pause and consider the extraordinary nature of knowledge synthesis in a sophisticated way, mirroring Perplexity's role in delivering clear answers from complex information.
The "real" component of this aesthetic is paramount for maintaining brand credibility and trust. Perplexity's core function is to provide "accurate, trusted, and real-time answers" with "citations". If the visuals were purely surreal or abstract, they could inadvertently undermine the brand's promise of verifiable and reliable knowledge. Consequently, the "real" aspect serves as an essential anchor, ensuring that the visuals remain relatable and grounded in a tangible reality, even when incorporating surreal elements. This suggests that AI models should prioritize photorealism or high-fidelity rendering as the base layer, upon which subtle surreal elements are layered, rather than starting from a highly stylized or abstract foundation.

1.3.2 Elements of "Retro Vibes" and "Editorial Modernity"

The "retro vibes" in Perplexity's branding are not about superficial nostalgia but are deeply intertwined with an "editorial and modern" aesthetic. This draws inspiration from "magazines and traditional media design," positioning Perplexity as a "trusted source of knowledge". This implies a "refined, curated look," with a focus on "composition, color balance," and a "sense of timelessness". General principles from "Retro Diffusion" regarding "intentional color limitations" and "proper techniques" for pixel art can be broadly applied, suggesting a disciplined approach to integrating retro elements.

The interpretation of "retro" as a signal of authority and curated knowledge is crucial. The explicit inspiration from "magazines and traditional media design" for the "editorial and modern" feel suggests that "retro vibes" are not about simple nostalgia. Instead, this aesthetic choice is intended to evoke the authority, trustworthiness, and curated nature of historical knowledge dissemination. Magazines, particularly well-designed ones from past eras, were often seen as reliable and meticulously presented sources of information. By subtly referencing this era, Perplexity's visuals can implicitly communicate a sense of reliability, depth, and thoughtful presentation, contrasting with the often fleeting or unverified nature of much digital information. This implies that AI prompts should prioritize classic compositional techniques, balanced layouts, and a refined color palette that feels timeless and professional, rather than overtly nostalgic or kitschy.

The consistent emphasis on "refined" and "modern" alongside "retro vibes" strongly suggests a restrained and sophisticated application of retro elements. This is not about overwhelming the viewer with obvious vintage effects or filters. Instead, it involves subtly integrating elements that evoke a classic, high-quality feel. The principle of "intentional color limitations" mentioned in the context of pixel art can be generalized: retro elements should be applied with precision and purpose, not as a blanket effect. This means AI prompts should specify subtle retro cues (e.g., "vintage color grading," "classic film grain," "editorial composition," "subtle chromatic aberration") rather than broad, aggressive "retro" styles that might detract from the "real" or "modern" aspects of the brand.

2. AI Image Generation Models: Capabilities for Branding Visuals

Selecting the appropriate AI image generation model is critical for consistently producing visuals aligned with Perplexity's unique brand aesthetic. This section evaluates leading models based on their specific strengths and weaknesses relevant to the "surreal-real with retro vibes" style.

2.1 GPT-Image-1: Strengths in Text Integration and Stylistic Range

GPT-Image-1, released in 2025 by OpenAI, represents a significant advancement in AI image generation, built upon the transformer architecture with enhanced vision and language capabilities. Its core strengths lie in precision text rendering, stylistic versatility, and advanced editing features.

The model excels at "Precision Text Rendering" in over 48 languages, naturally integrating text within images. This capability addresses a long-standing weakness in AI image generators, ensuring that text is not only visually correct but also contextually meaningful. For instance, it can generate road signs with proper warnings, book covers with coherent titles, and UI/UX

mockups with readable interface elements. This offers a strategic advantage for text-heavy brand communications. Perplexity's core offering is to "summarize information into concise answers with citations" , emphasizing clarity and precision. While not directly visual, this suggests that any brand visuals, particularly for marketing or product explanations, may require effective text integration. GPT-Image-1's precision text rendering directly addresses a historical weakness in AI image generation, making it uniquely suited for creating marketing collateral or educational content where readable, integrated text is paramount. This capability reduces the need for extensive post-processing and enhances the overall professionalism and clarity of brand visuals, aligning perfectly with Perplexity's "no gimmicks, no fluff" messaging.
GPT-Image-1 offers "Stylistic Versatility" with over 15 distinct visual approaches, including photorealistic, artistic, 3D rendering, graphic, and stylized options. It maintains high quality across these diverse styles, a notable advantage over many other AI generators. The model also features "Advanced Editing Features" like inpainting and outpainting, which maintain consistent lighting, perspective, and style across modified areas. The inclusion of an input_fidelity parameter allows granular control over how closely the generated output adheres to an input image's style, which is useful for "Maintaining brand identity". This capability is instrumental in bridging realism and stylization through fidelity control. GPT-Image-1's stylistic versatility encompasses both photorealistic and various artistic styles. The input_fidelity parameter allows granular control over how closely the generated output adheres to an input image's style and features. This is a critical capability for achieving the "surreal-real" aesthetic. One could initiate generation with a realistic base image (representing the "real") and then, through carefully crafted prompts and fidelity adjustments, introduce subtle surreal elements without sacrificing the grounded reality. This offers a powerful mechanism for controlled surrealism rather than random abstraction, enabling the precise blend required by Perplexity's sophisticated brand.

2.2 Midjourney v6: Advanced Realism and Prompt Coherence

Midjourney v6 introduces significant advancements, particularly in realism and prompt interpretation, making it highly relevant for Perplexity's brand visuals.
The model offers "more realistic and better overall quality," including "hyperrealism" that "does not diminish the model's creativity and stylistics," thereby allowing for "both realism and surrealism". Its "Text in Images" capability has also improved, overcoming previous limitations where text outputs were often messy or incomplete. This provides nuanced control for blending "surreal-real" elements. Midjourney V6's ability to generate hyperrealistic outputs while simultaneously supporting creativity and stylistics directly enables Perplexity's "surreal-real" aesthetic. This implies that the model can render realistic scenes with high fidelity and then, through careful prompting, introduce subtle or overt surreal elements without breaking the underlying realism. The "style tuner" and "enhanced prompt following" mean that complex instructions for blending these two elements (e.g., "a realistic office scene with a floating, glowing book") can be more accurately interpreted. This level of control is crucial for maintaining the brand's sophisticated, intentional surrealism rather than producing random, unguided outputs.
Midjourney v6 features "Enhanced Prompt Following" for longer (350+ words) and more complex natural language prompts, demonstrating improved coherence. For realistic outputs, the --style raw parameter is recommended instead of descriptive words like "photorealistic" or "8K". The ability to use detailed, natural language is important for conveying "retro vibes." Midjourney V6's shift towards understanding natural language and processing longer prompts is

a significant advantage. This allows for more descriptive and nuanced instructions when aiming for the "retro vibes" aesthetic. Instead of relying on simple keywords, a prompt engineer can describe the *feeling* of a retro image, the *texture* of old film, or the *compositional style* of vintage magazines. This aligns perfectly with Perplexity's "editorial and modern" interpretation of "retro". The model's ability to process such detailed descriptions means that subtle retro elements—like specific lighting, color grading, or compositional rules—can be precisely articulated in the prompt, leading to a refined, rather than kitschy, retro aesthetic.

2.3 FLUX.1: Character Consistency and Efficient Generation

FLUX.1, particularly its Kontext models, offers unique capabilities for maintaining visual consistency and achieving high-speed generation.
FLUX.1 provides "state-of-the-art character consistency" for preserving "unique elements... across multiple scenes and environments". It also supports "Local editing" for targeted modifications and "Style Reference" to generate new scenes while preserving unique styles from a reference image. This model boasts "Interactive Speed" with "inference speeds up to 8x faster than current leading models" like GPT-Image. However, it does have limitations, including occasional failures to follow instructions accurately and limited world knowledge, and the distillation process can introduce visual artifacts.
The model's strength in maintaining consistency can be leveraged for "brand motif consistency" for abstract concepts. While "character consistency" typically refers to human or animal characters, Perplexity's brand identity revolves around abstract concepts like "knowledge, learning, exploration," and "layered perspectives". FLUX.1's strength in maintaining consistency can be strategically leveraged to ensure that abstract visual motifs or symbolic representations of these concepts (e.g., a stylized "open book" element, the "intersecting windows" motif, or a recurring abstract shape representing information flow) remain visually consistent across diverse generated images. This is crucial for strengthening brand recognition and ensuring that the "surreal-real" elements, when recurring across different assets, appear cohesive and intentional rather than disjointed or random. This implies that FLUX.1 would be ideal for generating a series of visuals where a specific abstract element needs to be consistently rendered in various contexts.
FLUX.1's rapid inference speeds present a significant advantage for production environments requiring rapid iteration or high-volume asset generation. This speed must be balanced against the model's noted limitations in prompt following and its sensitivity to prompting-style. This implies a trade-off: while fast, FLUX.1 might require more precise or specialized prompting techniques to consistently achieve the desired "surreal-real with retro vibes" aesthetic. This suggests that for FLUX.1, the modular prompt system will need to be particularly robust in its core stylistic components, potentially relying less on complex narrative descriptions and more on structured keywords and parameters to guide the model effectively.

2.4 DALL-E 3: Strengths in Creativity and Limitations in Precision

DALL-E 3, developed by OpenAI, is known for its ability to create detailed and creative images from text prompts, with advanced language understanding and artistic flair.
DALL-E 3 excels at generating "highly detailed and creative images" and offers "diverse and often surprising results". It demonstrates strong artistic abilities, particularly with straightforward prompts. However, it struggles with "complex or abstract requests," sometimes misinterpreting instructions or failing to accurately depict requested scenes. For realism, generated images can

appear artificial, lacking natural imperfections. The model also has limitations in "accurate object representation" and "poor interpretation of abstract concepts".
The model's strengths lie in generating artistic interpretations for abstract concepts. DALL-E 3's capacity for creating "diverse and often surprising results" makes it suitable for exploring the "surreal" aspect of Perplexity's brand in a more unconstrained, creative manner. While it may struggle with precise control over complex compositions, its artistic flair can be leveraged for generating conceptual or abstract imagery that evokes the sense of wonder associated with knowledge discovery. This suggests DALL-E 3 is better suited for initial ideation or generating more abstract brand elements, where strict adherence to photorealism is less critical than artistic expression.
DALL-E 3's limitations in precision and consistency for complex prompts pose challenges for production-ready brand asset generation. The model's tendency to produce "artificial realism" and its difficulties with "accurate object representation" mean that achieving the "real" component of "surreal-real" with high fidelity may require significant iteration or manual refinement. Furthermore, its struggles with "complex or paradoxical prompts" indicate that detailed instructions for blending surreal and retro elements might be misinterpreted. This implies that for DALL-E 3, prompts should be simpler and more focused on abstract concepts or general artistic direction, with the understanding that post-generation editing might be necessary to align with precise brand guidelines. It is less suited for scenarios demanding strict visual reproducibility across multiple assets.

3. Modular, Production-Ready Prompt System for Adjacent Visuals

To reliably generate visuals consistent with Perplexity AI's "surreal-real with retro vibes" brand image across various AI models, a modular prompt system is essential. This approach enhances consistency, reusability, and control, making the process scalable for production environments.

3.1 Principles of Modular Prompt Engineering for Visual Consistency

Modular prompting involves structuring prompts into distinct segments, each targeting a specific task or behavior. This technique improves consistency, reusability, and control in large language model outputs by isolating context, instructions, examples, or goals into separate blocks. It represents a shift from single, massive "mega-prompts" to a library of reusable, interlocking "prompt modules". This method, often powered by HTML-like tags or JSON prompts, aims to transform AI from a flexible creative partner into a dependable engine for complex work.
This structured approach offers several advantages for brand visual identity. By breaking down prompts into smaller, focused modules, the AI model receives clearer instructions for each specific visual attribute, leading to more consistent and predictable results. Modules are self-contained prompt files with a clear purpose, making them reusable across different tasks and workflows, allowing for the creation of a library of components that can be easily integrated into various AI systems. Structured tags within modules create a "contract" with the AI, removing ambiguity and enforcing specific behaviors, tones, and objectives, thus providing immense control over the complex generation process. This architecture is robust, easy to maintain, and simple to upgrade, making it suitable for building scalable AI systems. Unlike brittle "mega-prompts," changes to one part of the process can be isolated to a specific module,

reducing the risk of unintended consequences. The use of structured tags transforms a complex prompt into a predictable, machine-readable function, which is the foundation of serious AI orchestration. It allows for breaking down massive tasks into a series of manageable, step-by-step prompting sequences, enabling multi-step generation with high accuracy.

3.2 Core Modules for Perplexity AI's Aesthetic

The following core modules are designed to capture Perplexity's "surreal-real with retro vibes" aesthetic. Each module can be combined or modified to generate a wide range of adjacent visuals.

3.2.1 Style and Aesthetic Module

This module defines the overarching visual style, blending realism, surrealism, and retro elements.

- **Keywords & Phrases:**
  - **Realism:** photorealistic, hyperrealistic, ultra-detailed, cinematic photography, DSLR quality, natural light, sharp focus.
  - **Surrealism (Subtle/Conceptual):** dreamlike atmosphere, subtle surreal elements, magical realism style, impossible juxtaposition, gravity-defying, ethereal, conceptual art, layered perspectives, symbolic imagery.
  - **Retro Vibes (Editorial Modernity):** vintage color grading, muted tones, soft film grain, classic editorial photography, mid-century modern aesthetic, analog feel, subtle chromatic aberration, timeless design.
  - **Overall Brand Tone:** intelligent and accessible, human-centric, curious, open, clean and confident.
- **Structure:** ,,,

3.2.2 Composition and Layout Module

This module dictates the arrangement of elements within the image, ensuring a structured yet inviting feel.

- **Keywords & Phrases:**
  - **Framing & Perspective:** wide shot, medium shot, close-up, eye-level perspective, low-angle, overhead view, rule of thirds, symmetrical composition, balanced composition.
  - **Arrangement:** minimalist composition, clean lines, uncluttered background, sparse elements, negative space emphasis, geometric arrangement, layered elements.
- **Structure:** ,

3.2.3 Color and Lighting Module

This module focuses on Perplexity's distinctive color palette and lighting to evoke trust and curiosity.

- **Keywords & Phrases:**
  - **Color Palette:** shades of teal, sky blue accents, warm earth tones, refined color palette, vibrant yet refined, monochromatic with teal highlights, analog color scheme.

    - **Lighting:** soft ambient lighting, diffused light, golden hour glow, natural window light, subtle shadows, volumetric lighting.
- **Structure:** ,

3.2.4 Subject and Concept Module

This module guides the generation of primary subjects, incorporating symbolic elements of knowledge and discovery.

- **Keywords & Phrases:**
    - **Core Concepts (Symbolic):** abstract representation of knowledge, interconnected data streams, floating geometric shapes, glowing information conduits, stylized open book, abstract cursor, expanding asterisk symbol.
    - **Human Element (Subtle):** silhouette of a curious figure, human hand interacting with digital interface, person observing a surreal landscape.
    - **Objects:** vintage technology elements, retro-futuristic devices, minimalist architecture.
- **Structure:** ``

3.2.5 Technical Specifications Module

This module ensures production-ready outputs by specifying technical parameters.

- **Keywords & Phrases:**
    - **Resolution:** 8K resolution, UHD, high definition.
    - **Aspect Ratio:** --ar 16:9 (for widescreen), --ar 3:2, --ar 1:1 (for square).
    - **Model-Specific Parameters:**
        - Midjourney: --v 6 (for v6 model), --style raw (for realism), --s [style_value] (for stylization strength), --seed [number] (for reproducibility).
        - GPT-Image-1: quality="high", style="photorealistic" or style="artistic", input_fidelity=[value].
        - DALL-E 3: quality=hd, style=vivid or style=natural, seed=[number] (within conversation).
        - FLUX.1: [specific FLUX.1 parameters for consistency] (Note: specific parameters not detailed in research, but concept of character consistency and style reference is present).
- **Structure:** ,,

3.2.6 Negative Prompt Module

This module defines elements and qualities to explicitly avoid, ensuring a clean and consistent output.

- **Keywords & Phrases:** no distortions, no blurry elements, no chaotic composition, no oversaturated colors, no cartoonish elements (unless specified in style module), no generic tech tropes.
- **Structure:** --no (Midjourney syntax) or Avoid (general)

3.3 Model-Specific Adaptations and Parameters

While the core modules provide a universal framework, optimal results require model-specific

adaptations.

- **GPT-Image-1:**
  - **Strengths leveraged:** Precision text rendering and broad stylistic versatility. The input_fidelity parameter is crucial for blending "real" base images with subtle "surreal" elements.
  - **Prompting considerations:** Utilize detailed descriptions for text elements. Experiment with input_fidelity to control the degree of surrealism while preserving realism.
  - **Example:** A photorealistic library scene, with books subtly floating towards a central glowing orb, vintage color grading, clean lines. Text on a floating book: "Where Knowledge Begins". --quality=high --style=photorealistic --input_fidelity=0.7
- **Midjourney v6:**
  - **Strengths leveraged:** Advanced realism, enhanced prompt following for natural language, and ability to blend realism and stylistics.
  - **Prompting considerations:** Employ the --style raw parameter for the "real" component. Use descriptive, natural language for "retro vibes" and subtle surreal elements. Leverage its improved coherence for complex multi-object prompts.
  - **Example:** A hyperrealistic, serene landscape with a large, transparent geometric structure gently hovering above a tranquil lake, reflecting a muted sunset. The scene has a subtle vintage film grain and a refined, curated aesthetic reminiscent of classic magazine photography. --ar 16:9 --v 6 --style raw
- **FLUX.1:**
  - **Strengths leveraged:** State-of-the-art character/element consistency and high inference speed. This is valuable for maintaining consistent symbolic motifs across a series of images.
  - **Prompting considerations:** Focus on clear, structured prompts for core elements that need to be consistent. Given its occasional prompt following limitations, simpler, more direct instructions for complex surreal elements might be necessary, or iterative refinement may be required.
  - **Example:** A photorealistic image of a stylized open book, glowing with soft teal light, casting a subtle shadow on a warm earth-toned surface. The book's pages are subtly rippling as if filled with unseen data streams. Minimalist composition, editorial lighting. --seed 12345 (for consistency)
- **DALL-E 3:**
  - **Strengths leveraged:** Strong artistic abilities and diverse, creative outputs for more abstract concepts.
  - **Prompting considerations:** Best for exploring the "surreal" aspect with less emphasis on strict photorealism or precise object placement. Prompts should be detailed but allow for creative interpretation. Leverage its strength in combining multiple inputs for complex conceptual imagery.
  - **Example:** An abstract, dreamlike composition featuring interconnected teal and sky blue geometric shapes floating above a warm, textured ground, evoking a sense of discovery and knowledge flow. The scene has a subtle retro-futuristic feel, like a vintage magazine illustration. --quality=hd --style=vivid

3.4 Workflow for Production-Ready Visual Generation

A structured workflow is critical for ensuring reliability and reproducibility in a production

environment.

1. **Concept Definition:** Clearly define the visual concept for the asset, linking it explicitly to Perplexity's brand philosophy and the "surreal-real with retro vibes" aesthetic.
2. **Modular Prompt Assembly:** Assemble the prompt using the core modules (Style, Composition, Color/Lighting, Subject/Concept, Technical Specs, Negative Prompt). Start with a general prompt and progressively add detail.
3. **Model Selection & Initial Generation:** Choose the most suitable AI model based on the specific requirements of the visual (e.g., GPT-Image-1 for text integration, Midjourney for nuanced realism). Generate initial iterations.
4. **Iterative Refinement:**
    - **Analyze Output:** Evaluate generated images against brand guidelines and the desired aesthetic.
    - **Adjust Modules:** Modify specific modules (e.g., tweak color values, adjust surreal elements, refine composition keywords) to guide the AI towards the target.
    - **Leverage Model-Specific Parameters:** Utilize parameters like --seed (Midjourney, DALL-E 3) or input_fidelity (GPT-Image-1) to maintain consistency across variations and control stylistic adherence.
    - **Negative Prompts:** Continuously refine negative prompts to eliminate undesired elements or visual artifacts.
5. **Version Control and Documentation:**
    - **Prompt Versioning:** Maintain a version-controlled repository of all successful prompts and their variations. This includes the full prompt string, the AI model used, and any specific parameters or seeds.
    - **Output Archiving:** Archive generated images alongside their corresponding prompts and metadata for future reference and reproducibility.
    - **Feedback Loop:** Document feedback on generated visuals and how prompt adjustments addressed them, creating a learning database for future generations.
6. **Quality Assurance:** Conduct a final review to ensure the visual assets meet all brand standards and technical specifications before deployment. Reproducibility in machine learning, which includes AI image generation, means that algorithms can be repeatedly run on specific datasets to obtain the same or similar results. This is crucial for scalability and production deployment. Tracking changes in code, data, and environment is essential for achieving this.

4. Conclusions and Recommendations

Perplexity AI's brand image, characterized by a "surreal-real with retro vibes," is a sophisticated and deliberate choice designed to differentiate it in the AI market by projecting intelligence, accessibility, and a human-centric approach to knowledge. The "surreal-real" aspect visually communicates the profound capabilities of AI in synthesizing information while maintaining a grounded, trustworthy reality. The "retro vibes," interpreted as "editorial modernity," evoke a sense of curated knowledge and timeless reliability, contrasting with the fleeting nature of much digital content.

To reliably generate adjacent visuals across top AI image models, a modular prompt system is not merely advantageous but essential. This system allows for precise control over stylistic elements, compositional structure, color palettes, and symbolic content, ensuring that generated visuals consistently align with Perplexity's unique brand identity. GPT-Image-1 excels in text

integration and fidelity control, making it ideal for visuals requiring integrated text and controlled blending of realism and stylization. Midjourney v6 offers advanced realism and superior natural language understanding, allowing for nuanced descriptions of complex scenes and subtle retro aesthetics. While FLUX.1 provides high-speed generation and strong consistency for recurring abstract motifs, its prompt adherence requires careful management. DALL-E 3, with its strong artistic capabilities, is best suited for exploring more abstract or conceptual surreal elements, though it may require more post-processing for precise realism.

Actionable Recommendations:

1. **Implement a Centralized Modular Prompt Library:** Develop and maintain a version-controlled library of prompt modules, categorized by aesthetic components (Style, Composition, Color/Lighting, Subject/Concept, Technical, Negative). This ensures consistency and reusability across teams and projects.
2. **Establish Model-Specific Prompting Guidelines:** Create detailed guidelines for adapting the core modules to each AI model (GPT-Image-1, Midjourney v6, FLUX.1, DALL-E 3), accounting for their unique strengths, weaknesses, and parameter syntaxes. This will optimize output quality and efficiency for each tool.
3. **Prioritize Iterative Refinement with Seed Management:** Adopt an iterative workflow where initial generations are systematically refined. For models supporting it, consistently use and document seed numbers to ensure reproducibility, especially for core brand elements or recurring motifs.
4. **Integrate Negative Prompting as a Standard Practice:** Actively utilize negative prompts within each module to precisely sculpt the desired aesthetic by eliminating undesirable visual artifacts or stylistic deviations.
5. **Conduct Regular Brand Visual Audits:** Periodically review AI-generated assets against established brand guidelines to ensure continued fidelity to the "surreal-real with retro vibes" aesthetic. This helps identify any drift in the AI models' interpretations over time or with new model updates.
6. **Invest in Cross-Functional Training:** Provide training for creative, marketing, and technical teams on the modular prompt system and model-specific nuances. This fosters a shared understanding of the brand's visual language and empowers more users to generate on-brand assets.

By adopting this structured and adaptable approach to AI image generation, Perplexity AI can reliably produce visual assets that not only resonate with its distinctive brand image but also scale efficiently for diverse marketing and product needs, reinforcing its unique position as "where knowledge begins."

Works cited

1. Brand Like the Best: The Perplexity Brand Identity - Numinous®, https://numinousco.com/the-perplexity-brand-identity/ 2. About Perplexity, https://www.perplexity.ai/hub/about 3. Perplexity Brand Identity: Design & Strategy Insights - BytePlus, https://www.byteplus.com/en/topic/419690 4. AI and Magical Realism: When technology blurs the line between wonder and reality - Diplo, https://www.diplomacy.edu/blog/ai-and-magical-realism-when-technology-blurs-the-line-between-wonder-and-reality/ 5. A Study on the Expression Techniques of Surrealism in Realistic Media Art: focusing on d'strict's works - ResearchGate, https://www.researchgate.net/publication/374330694_A_Study_on_the_Expression_Techniques_of_Surrealism_in_Realistic_Media_Art_focusing_on_d'strict's_works 6. Perplexity AI,

https://www.perplexity.ai/ 7. What is Perplexity AI? A Smarter Way to Search | DigitalOcean, https://www.digitalocean.com/resources/articles/what-is-perplexity-ai 8. Retro Diffusion: Creating authentic pixel art with AI at scale | Runware, https://runware.ai/blog/retro-diffusion-creating-authentic-pixel-art-with-ai-at-scale 9. GPT-Image-1: OpenAI Image Generator Model and Its Changing ..., https://mpgone.com/gpt-image-1-openai-image-generator-model-and-its-changing-effect/ 10. Perplexity in Practice: Putting AI to work for marketing teams, https://www.perplexity.ai/enterprise/pplx-in-practice-marketing-replay 11. What's new in Azure OpenAI in Azure AI Foundry Models? | Microsoft Learn, https://learn.microsoft.com/en-us/azure/ai-foundry/openai/whats-new 12. Midjourney V6 Review (What's New?) - TextCortex, https://textcortex.com/post/midjourney-v6-review 13. 9 Midjourney V6. Prompting Technique You Need to Know | by Christie C. - Medium, https://medium.com/design-bootcamp/9-midjourney-v6-prompting-technique-you-need-to-know-b467aa07365e 14. Prompt Basics - Midjourney, https://docs.midjourney.com/hc/en-us/articles/32023408776205-Prompt-Basics 15. Introducing FLUX.1 Kontext and the BFL Playground | Black Forest Labs, https://bfl.ai/announcements/flux-1-kontext 16. black-forest-labs/FLUX.1-dev - Hugging Face, https://huggingface.co/black-forest-labs/FLUX.1-dev 17. DALL-E 3: An In-Depth AI Image Generator Review | FlowHunt, https://www.flowhunt.io/blog/dall-e-3-an-in-depth-ai-image-generator-review/ 18. Modular Prompting - Prompt Engineering For Scale, https://optizenapp.com/ai-prompts/modular-prompting/ 19. Advanced Prompt Techniques: Getting Hyper-Realistic Results from ..., https://stockimg.ai/blog/prompts/advanced-prompt-techniques-getting-hyper-realistic-results-from-your-ai-photo-generator 20. The Best 25 Midjourney Prompts for Minimalist - OpenArt, https://openart.ai/blog/post/midjourney-prompts-for-minimalist 21. Retro-futurist Image Prompts • PromptDen, https://promptden.com/inspiration/retro-futurist+all 22. Retro-futuristic Image Prompts - PromptDen, https://promptden.com/inspiration/retro-futuristic+all 23. What Does SEED Mean in AI Art | Everything You Need to Know ..., https://shaicreative.ai/what-does-seed-mean-in-ai-art-everything-you-need-to-know/ 24. GPT Image 1: OpenAI's Powerful Image Generation Model - Peerlist, https://peerlist.io/blog/engineering/gpt-image-1-powerful-image-generation-model-by-openai 25. Ultimate DALL·E Prompt Guide: Unlocking The Secrets Of Effective DALL-E Prompts For Stunning Creations - Brain Pod AI, https://brainpod.ai/ultimate-dall%C2%B7e-prompt-guide-unlocking-the-secrets-of-effective-dall-e-prompts-for-stunning-creations/ 26. Consistent Image generation for Story using DALLE - API - OpenAI Developer Community, https://community.openai.com/t/consistent-image-generation-for-story-using-dalle/612276 27. How to Solve Reproducibility in ML - neptune.ai, https://neptune.ai/blog/how-to-solve-reproducibility-in-ml